

# Convergence diagnostics for MCMC chains

MLPM Summer School 2015

Catalina A. Vallejos, Nils Eling, John C. Marioni

22 September 2015

# Convergence diagnostics for MCMC chains

To assess convergence is not a trivial task, especially in large-dimensional settings . . .

What can we do?

- ▶ Visual inspection: trace-plots, cumulative means plots, etc
- ▶ Statistical tests: e.g. the ones implemented in the coda library

We can also run multiple chains (with different starting values) and check if all chains converge to the same values.

# Convergence diagnostics for MCMC chains

To illustrate convergence diagnostic tools, we create 4 **artificial** chains

```
set.seed(1)
Chain1 = rnorm(n = 10000, mean = 5, sd = 2)
set.seed(2)
Chain2 = c(rnorm(n = 5000, mean = (1:5000)/1000, sd = 1),
           rnorm(n = 5000, mean = 5, sd = 1))
set.seed(3)
Chain3 = c(rep(1, 2000), rep(-2, 500),
           rep(2, 5000), rep(5, 2500)) +
           rnorm(n = 10000, mean = 0, sd = 1)
set.seed(4)
Chain4 = sin((1:10000)/1000) + rnorm(10000, 0, 1.5)
```

# Convergence diagnostics for MCMC chains


## Visual inspection - Traceplots

```
par(mfrow = c(2,2))
plot(Chain1, type = "l", bty = "n", xlab = "Iterations",
     ylab = "Draws", main = expression("Good mixing"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
plot(Chain2, type = "l", bty = "n", xlab = "Iterations",
     ylab = "Draws", main = expression("Additional burning required"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
plot(Chain3, type = "l", bty = "n", xlab = "Iterations",
     ylab = "Draws", main = expression("Stuck at local modes"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
plot(Chain4, type = "l", bty = "n", xlab = "Iterations",
     ylab = "Draws", main = expression("Oscilating"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
```


# Convergence diagnostics for MCMC chains

## Visual inspection - Traceplots


Good mixing




Additional burning required



Stuck at local modes



Oscilating



# Convergence diagnostics for MCMC chains


## Visual inspection - Cumulative means plots

```
par(mfrow = c(2,2))
plot(cumsum(Chain1)/1:10000, type = "l", bty = "n",
     xlab = "Iterations", ylab = "Cumulative mean",
     main = expression("Good mixing"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
plot(cumsum(Chain2)/1:10000, type = "l", bty = "n",
     xlab = "Iterations", ylab = "Cumulative mean",
     main = expression("Additional burning required"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
plot(cumsum(Chain3)/1:10000, type = "l", bty = "n",
     xlab = "Iterations", ylab = "Cumulative mean",
     main = expression("Stuck at local modes"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
plot(cumsum(Chain4)/1:10000, type = "l", bty = "n",
     xlab = "Iterations", ylab = "Draws",
     main = expression("Oscilating"),
     cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
```


# Convergence diagnostics for MCMC chains

## Visual inspection - Cumulative means plots


Good mixing




Additional burning required



Stuck at local modes



Oscilating



# Convergence diagnostics for MCMC chains





## Visual inspection - Auto-correlation plots

```
par(mfrow = c(2,2))
acf(Chain1, bty = "n",
    main = expression("Good mixing"),
    cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
acf(Chain2, bty = "n",
    main = expression("Additional burning required"),
    cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
acf(Chain3, bty = "n",
    main = expression("Stuck at local modes"),
    cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
acf(Chain4, bty = "n",
    main = expression("Oscilating"),
    cex.main = 2.5, cex.lab = 1.5, cex.axis = 1.5)
```



# Convergence diagnostics for MCMC chains

## Visual inspection - Auto-correlation plots



# Convergence diagnostics for MCMC chains

## Statistical tests

The R library makes available several statistical test for convergence diagnostics. To load this library use

```
library(coda)
```

To use these tools, a first step is to transform chains into objects of class `mcmc`.

```
Chain1.mcmc = mcmc(Chain1)
```

```
Chain2.mcmc = mcmc(Chain2)
```

```
Chain3.mcmc = mcmc(Chain3)
```

```
Chain4.mcmc = mcmc(Chain4)
```

# Convergence diagnostics for MCMC chains

## Statistical tests - Geweke

Compares the means of two non-overlapping segments of the chain (typically the first 10% and last 50% of draws). It returns a z-score which is adjusted for autocorrelation

```
geweke.diag(Chain1.mcmc)$z
```

```
##          var1  
## -0.02164115
```

```
geweke.diag(Chain2.mcmc)$z
```

```
##          var1  
## -57.39918
```

```
geweke.diag(Chain3.mcmc)$z
```

```
##          var1  
## -2.722006
```

```
geweke.diag(Chain4.mcmc)$z
```

```
##          var1  
## 1.416636
```

# Convergence diagnostics for MCMC chains

## Statistical tests

Other diagnostics implemented in coda include

- ▶ Gelman and Rubin
- ▶ Raftery and Lewis
- ▶ Heidelberg and Welch

Because of time restrictions, we won't explore all of these today.

NOTE: For better results it is important to combine visual inspection and test-based convergence diagnostics

# Convergence diagnostics for MCMC chains

Each of these convergence diagnostics are designed to assess the convergence of a chain for a single parameter

*What to do for high-dimensional models?*

Once again, there is no trivial answer. Some ideas:

- ▶ Select a random subset of model parameters
- ▶ Calculate averages within sets of parameters